

# CS/DS541 Deep Learning

Joshua Levy  
Sreejani Chatterjee

## Part1:

Trained for 10 Epochs, and batch size of 64

HOMEWORK5_JLEVY_SCHATTERJEE	1	Model: "sequential"																																	
homework5_jlevy_schatterjee.ipynb	2																																		
	3																																		
	4	<table><thead><tr><th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr></thead><tbody><tr><td>conv2d (Conv2D)</td><td>(None, 28, 28, 64)</td><td>320</td></tr><tr><td>max_pooling2d (MaxPooling2D)</td><td>(None, 27, 27, 64)</td><td>0</td></tr><tr><td>dropout (Dropout)</td><td>(None, 27, 27, 64)</td><td>0</td></tr><tr><td>conv2d_1 (Conv2D)</td><td>(None, 27, 27, 32)</td><td>8224</td></tr><tr><td>max_pooling2d_1 (MaxPooling2D)</td><td>(None, 26, 26, 32)</td><td>0</td></tr><tr><td>dropout_1 (Dropout)</td><td>(None, 26, 26, 32)</td><td>0</td></tr><tr><td>flatten (Flatten)</td><td>(None, 21632)</td><td>0</td></tr><tr><td>dense (Dense)</td><td>(None, 256)</td><td>5538048</td></tr><tr><td>dropout_2 (Dropout)</td><td>(None, 256)</td><td>0</td></tr><tr><td>dense_1 (Dense)</td><td>(None, 10)</td><td>2570</td></tr></tbody></table>	Layer (type)	Output Shape	Param #	conv2d (Conv2D)	(None, 28, 28, 64)	320	max_pooling2d (MaxPooling2D)	(None, 27, 27, 64)	0	dropout (Dropout)	(None, 27, 27, 64)	0	conv2d_1 (Conv2D)	(None, 27, 27, 32)	8224	max_pooling2d_1 (MaxPooling2D)	(None, 26, 26, 32)	0	dropout_1 (Dropout)	(None, 26, 26, 32)	0	flatten (Flatten)	(None, 21632)	0	dense (Dense)	(None, 256)	5538048	dropout_2 (Dropout)	(None, 256)	0	dense_1 (Dense)	(None, 10)	2570
Layer (type)	Output Shape	Param #																																	
conv2d (Conv2D)	(None, 28, 28, 64)	320																																	
max_pooling2d (MaxPooling2D)	(None, 27, 27, 64)	0																																	
dropout (Dropout)	(None, 27, 27, 64)	0																																	
conv2d_1 (Conv2D)	(None, 27, 27, 32)	8224																																	
max_pooling2d_1 (MaxPooling2D)	(None, 26, 26, 32)	0																																	
dropout_1 (Dropout)	(None, 26, 26, 32)	0																																	
flatten (Flatten)	(None, 21632)	0																																	
dense (Dense)	(None, 256)	5538048																																	
dropout_2 (Dropout)	(None, 256)	0																																	
dense_1 (Dense)	(None, 10)	2570																																	
	5																																		
	6																																		
	7																																		
	8																																		
	9																																		
	10																																		
	11																																		
	12																																		
	13																																		
	14																																		
	15																																		
	16																																		
	17																																		
	18																																		
	19																																		
	20																																		
	21																																		
	22																																		
	23																																		
	24																																		
	25																																		
	26																																		
	27																																		
	28	Total params: 5,549,162																																	
	29	Trainable params: 5,549,162																																	
	30	Non-trainable params: 0																																	
	31																																		
	32																																		

## Architecture used:

Convolution layer with 64 filters, each 2x2, stride of 1, without padding activation relu

Max pool with a pooling width of 2x2

Dropout of 0.3, Dropout layer randomly sets input units to 0 with a frequency of rate at each step during training time, which helps prevent overfitting

Convolution layer with 32 filters, each 2x2, stride of 1, without padding activation relu

Max pool with a pooling width of 2x2

Dropout of 0.3,

Flattened all feature maps into one long vector.

Fully-connected layer to map into a 256-dimensional vector

Relu

Dropout of 0.5

Fully-connected layer to map into a 10-dimensional vector.

Softmax.

## Results:

Loss and accuracy values over the epochs

```
HOMEWORKS_ILEVY_SCHATTERJEE
homework5_ilevy_schatterjee.ipynb
model.weights.best.hdf5

1 860/860 [=====] - ETA: 0s - loss: 0.4893 - accuracy: 0.8236
2 Epoch 1: val_loss improved from inf to 0.31657, saving model to model.weights.best.hdf5
3 860/860 [=====] - 31s 33ms/step - loss: 0.4893 - accuracy: 0.8236 - val_loss: 0.3166 - val_accuracy: 0.8868
4 Epoch 2/10
5 859/860 [=====>.] - ETA: 0s - loss: 0.3382 - accuracy: 0.8788
6 Epoch 2: val_loss improved from 0.31657 to 0.27397, saving model to model.weights.best.hdf5
7 860/860 [=====] - 29s 33ms/step - loss: 0.3381 - accuracy: 0.8788 - val_loss: 0.2740 - val_accuracy: 0.9018
8 Epoch 3/10
9 859/860 [=====>.] - ETA: 0s - loss: 0.2919 - accuracy: 0.8917
10 Epoch 3: val_loss improved from 0.27397 to 0.25273, saving model to model.weights.best.hdf5
11 860/860 [=====] - 29s 33ms/step - loss: 0.2919 - accuracy: 0.8917 - val_loss: 0.2527 - val_accuracy: 0.9110
12 Epoch 4/10
13 859/860 [=====>.] - ETA: 0s - loss: 0.2615 - accuracy: 0.9042
14 Epoch 4: val_loss improved from 0.25273 to 0.23051, saving model to model.weights.best.hdf5
15 860/860 [=====] - 29s 33ms/step - loss: 0.2618 - accuracy: 0.9041 - val_loss: 0.2305 - val_accuracy: 0.9154
16 Epoch 5/10
17 859/860 [=====>.] - ETA: 0s - loss: 0.2371 - accuracy: 0.9123
18 Epoch 5: val_loss improved from 0.23051 to 0.20376, saving model to model.weights.best.hdf5
19 860/860 [=====] - 29s 33ms/step - loss: 0.2371 - accuracy: 0.9123 - val_loss: 0.2038 - val_accuracy: 0.9264
20 Epoch 6/10
21 859/860 [=====>.] - ETA: 0s - loss: 0.2199 - accuracy: 0.9182
22 Epoch 6: val_loss improved from 0.20376 to 0.19705, saving model to model.weights.best.hdf5
23 860/860 [=====] - 29s 34ms/step - loss: 0.2200 - accuracy: 0.9181 - val_loss: 0.1970 - val_accuracy: 0.9254
24 Epoch 7/10
25 859/860 [=====>.] - ETA: 0s - loss: 0.2020 - accuracy: 0.9254
26 Epoch 7: val_loss did not improve from 0.19705
27 860/860 [=====] - 29s 34ms/step - loss: 0.2020 - accuracy: 0.9254 - val_loss: 0.1991 - val_accuracy: 0.9274
28 Epoch 8/10
29 859/860 [=====>.] - ETA: 0s - loss: 0.1858 - accuracy: 0.9316
30 Epoch 8: val_loss did not improve from 0.19705
31 860/860 [=====] - 29s 34ms/step - loss: 0.1858 - accuracy: 0.9316 - val_loss: 0.2028 - val_accuracy: 0.9240
32 Epoch 9/10
33 859/860 [=====>.] - ETA: 0s - loss: 0.1761 - accuracy: 0.9334
34 Epoch 9: val_loss improved from 0.19705 to 0.19180, saving model to model.weights.best.hdf5
35 860/860 [=====] - 29s 34ms/step - loss: 0.1760 - accuracy: 0.9334 - val_loss: 0.1918 - val_accuracy: 0.9324
36 Epoch 10/10
37 859/860 [=====>.] - ETA: 0s - loss: 0.1667 - accuracy: 0.9374
38 Epoch 10: val_loss improved from 0.19180 to 0.18758, saving model to model.weights.best.hdf5
39 860/860 [=====] - 29s 34ms/step - loss: 0.1667 - accuracy: 0.9374 - val_loss: 0.1876 - val_accuracy: 0.9354
40
```

```
▶ #Load model with best validation accuracy
model.load_weights("model.weights.best.hdf5")
#Test accuracy
score = model.evaluate(xTest, yTest, verbose = 0)
print('\n', 'Test accuracy: ', score[1])

[34] ✓ 0.5s

...

Test accuracy: 0.9279999732971191
```

**Test Accuracy Obtained: 92.79%**

## Part2:

Trained for 5 Epochs, and batch size of 64

### Architecture used:

Convolution layer with 64 filters, each 3x3, stride of 1 no padding.

Max pool with a pooling width of 2x2, stride of 2, no padding.

ReLU.

Flatten the 64 feature maps into one long vector.

Fully-connected layer to map into a 1024-dimensional vector.

ReLU.

Fully-connected layer to map into a 10-dimensional vector.

Softmax.

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_6 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_6 (MaxPooling 2D)	(None, 13, 13, 64)	0
re_lu (ReLU)	(None, 13, 13, 64)	0
flatten_3 (Flatten)	(None, 10816)	0
dense_6 (Dense)	(None, 1024)	11076608
re_lu_1 (ReLU)	(None, 1024)	0
dense_7 (Dense)	(None, 10)	10250
softmax (Softmax)	(None, 10)	0

=====

Total params: 11,087,498  
Trainable params: 11,087,498  
Non-trainable params: 0

## Results:

### Loss and accuracy values over the epochs

```
#Compile model
model_p2.compile(loss = 'categorical_crossentropy', optimizer = 'adam', metrics = ['accuracy'])

from keras.callbacks import ModelCheckpoint
checkpointer = ModelCheckpoint(filepath = 'model_p2.weights.best.h5', verbose = 1, save_best_only = True)
#Train the model
model_p2.fit(xTrain, yTrain, batch_size = 64, epochs = 5, validation_data = (xValid, yValid), callbacks = [checkpointer])

[40] ✓ 12.7s

...
Epoch 1/5
852/860 [=====>.] - ETA: 0s - loss: 0.0353 - accuracy: 0.9875
Epoch 1: val_loss improved from inf to 0.41162, saving model to model_p2.weights.best.h5

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op while saving (showing 1 of 1). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: model_p2.weights.best.h5/assets
INFO:tensorflow:Assets written to: model_p2.weights.best.h5/assets

860/860 [=====] - 3s 4ms/step - loss: 0.0351 - accuracy: 0.9875 - val_loss: 0.4116 - val_accuracy: 0.9192
Epoch 2/5
859/860 [=====>.] - ETA: 0s - loss: 0.0221 - accuracy: 0.9932
Epoch 2: val_loss did not improve from 0.41162
860/860 [=====] - 2s 3ms/step - loss: 0.0221 - accuracy: 0.9932 - val_loss: 0.4152 - val_accuracy: 0.9246
Epoch 3/5
857/860 [=====>.] - ETA: 0s - loss: 0.0158 - accuracy: 0.9949
Epoch 3: val_loss did not improve from 0.41162
860/860 [=====] - 2s 3ms/step - loss: 0.0159 - accuracy: 0.9949 - val_loss: 0.4592 - val_accuracy: 0.9222
Epoch 4/5
857/860 [=====>.] - ETA: 0s - loss: 0.0205 - accuracy: 0.9929
Epoch 4: val_loss did not improve from 0.41162
860/860 [=====] - 2s 3ms/step - loss: 0.0206 - accuracy: 0.9929 - val_loss: 0.4818 - val_accuracy: 0.9168
Epoch 5/5
859/860 [=====>.] - ETA: 0s - loss: 0.0111 - accuracy: 0.9963
Epoch 5: val_loss did not improve from 0.41162
860/860 [=====] - 2s 3ms/step - loss: 0.0111 - accuracy: 0.9963 - val_loss: 0.5166 - val_accuracy: 0.9148

<keras.callbacks.History at 0x7f33cb8d7b50>
```

```
#Test accuracy
score = model_p2.evaluate(xTest, yTest, verbose = 0)
print('\n', 'Test accuracy: ', score[1])

90] ✓ 0.3s

**

Test accuracy: 0.9161999821662903
```

### Test Accuracy Obtained: 91.62%

```
This is from Tensorflow Model [3.3310093e-12 1.9235139e-19 5.0639418e-14 1.6371341e-14 7.1424300e-15
1.6809189e-11 1.6845696e-13 1.0000000e+00 1.4049499e-12 9.4282771e-15]
This is from Numpy Model [3.33100924e-12 1.92351659e-19 5.06393396e-14 1.63713661e-14
7.14242818e-15 1.68091697e-11 1.68456940e-13 1.00000000e+00
1.40494739e-12 9.42827208e-15]
This value shows similarity between two layers 100.0
```

## Comparison of results from Tensorflow and Numpy code:

We verify the accuracy by predicting a random test data for both the Tensorflow and Numpy model.

### Tensorflow:

```
# predicting 1 random sample test data to compare later with numpy forward propagation
yhat1 = model_p2.predict(xTest[998:999,:,:,:])[0]
print(yhat1)
```

[53] ✓ 0.6s

... 1/1 [=====] - 0s 15ms/step  
[3.3310093e-12 1.9235139e-19 5.0639418e-14 1.6371341e-14 7.1424300e-15  
1.6809189e-11 1.6845696e-13 1.0000000e+00 1.4049499e-12 9.4282771e-15]

### Numpy:

```
# Working on the same test data xTest[998]

yhat = softmax(h3)
#print(yhat.Shape)
yhat2 = yhat
yhat2 = yhat.reshape(-1)
print(yhat2)
```

[64] ✓ 0.1s

... [3.33100924e-12 1.92351659e-19 5.06393396e-14 1.63713661e-14  
7.14242818e-15 1.68091697e-11 1.68456940e-13 1.00000000e+00  
1.40494739e-12 9.42827208e-15]

```
This is from Tensorflow [3.3310093e-12 1.9235139e-19 5.0639418e-14 1.6371341e-14 7.1424300e-15  
1.6809189e-11 1.6845696e-13 1.0000000e+00 1.4049499e-12 9.4282771e-15]  
This is from Network through Numpy [3.33100924e-12 1.92351659e-19 5.06393396e-14 1.63713661e-14  
7.14242818e-15 1.68091697e-11 1.68456940e-13 1.00000000e+00  
1.40494739e-12 9.42827208e-15]  
This is accuracy which shows how two layers are similar 100.0
```

### Results:

Both outputs have same values leading to a 100% similarity in results

The following plot shows the same.

